

TONY WANG

US CITIZEN | 425-221-4924 | wangtzero@gmail.com | linkedin.com/in/tony-wang-66667b242

EDUCATION

Georgia Institute of Technology
Bachelor of Science in Computer Science
Relevant Coursework: Data Structures & Algorithms, Machine Learning, Computer Organization and Programming, Discrete Math, Linear Algebra, Object-Oriented Programming (Java)

May 2027
Atlanta, GA

TECHNICAL SKILLS

Languages: Python, C, C++, Java, JavaScript, SQL, Bash, Prompt-Engineering
Systems/Performance: Linux, Cisco UCS, CUDA, NUMA, NCCL, NVLink, PCIe, Intel MLC, SPEC CPU, perf
AI/ML Infrastructure: PyTorch, TensorFlow, vLLM, llama.cpp, Kubernetes, Docker, Prometheus, Grafana
Tools & Frameworks: FastAPI, Redis, REST APIs, Git, GitHub, JIRA, React.js, Node.js

EXPERIENCE

Cisco
Software Engineer Intern

- Optimized AI server infrastructure through benchmarking, observability, and performance analysis on Cisco UCS systems
- Compared next-gen Intel CPU and NVIDIA GPU ML inference, analyzing latency, throughput, memory, and tradeoffs
- Reduced inference profiling time by **15%** by automating capture of KV cache, batch size, thread scaling, and telemetry data
- Built Kubernetes LLM benchmarking workflows, comparing multi-GPU inference scaling across throughput and latency
- Enabled real-time server observability by deploying **Prometheus/Grafana** dashboards for health and workload monitoring

May 2026 – Present
San Jose, CA

Lenovo
Systems Software Engineer Intern

- Partnered with NVIDIA CUDA engineers to debug NCCL bottlenecks, improving GPU cluster efficiency by **15%** overall
- Analyzed NVLink/PCIe conflicts with HPC diagnostics, guiding targeted upgrades that removed hardware bottlenecks
- Diagnosed InfiniBand/NCCL latency with ibstat and NCCL traces, reducing distributed GPU workload runtime by **20%**
- Reduced recurring distributed workload failures by debugging Linux routing and DNS issues across GPU cluster nodes
- Automated Docker CI/CD validation workflows with Python, reducing manual testing time by **40%** overall consistently

Jun 2024 – Aug 2024
Morrisville, NC

Hong Kong University of Science and Technology
Machine Learning Engineer Intern

- Engineered Python pipelines processing **1,000+** patient records, preserving **99%** data integrity for medical AI experiments
- Optimized Random Forest regressors with cross-validation, reducing overfitting and reaching a **0.79 R²** validation score
- Refactored SQL/Bash ETL workflows for clinical datasets, reducing ingestion latency by **30%** across preprocessing runs
- Parallelized cross-validation across **256** drug combinations, accelerating ranking workflows for clinical research teams

Jun 2025 – Aug 2025
Guangzhou, China

PROJECTS

Personal Website | *React, Vite, Vercel, GitHub, Codex*

- Built responsive technical portfolio with React and Vite to showcase experience, projects, education, and personal timeline
- Used Codex and prompt-engineering workflows to iterate on layout, content structure, polish, and responsive behavior
- Designed custom CSS visuals and recruiter-facing sections, deploying version-controlled releases through Vercel

Live | GitHub

Linux Performance Profiler | *C, Python, Bash, Linux, Intel MLC*

- Implemented C microbenchmarks and Bash runners measuring memory, storage, and process latency across **3** workloads
- Generated Python reports for throughput, **p50/p95** latency, and run variance across **30+** benchmark trials
- Added CSV logging and threshold checks to flag **10%+** performance regressions across repeated Linux benchmark runs
- Compared CPU, memory, and NVMe bottlenecks on Linux, documenting performance tradeoffs across **3** system paths

The Dining Council | *FastAPI, Redis, Docker, Streamlit, Gemini 2.5, Google Maps API*

- Developed AI-powered group restaurant recommendation app with shared sessions, live chat, and **6**-character room codes
- Built FastAPI backend with **10+** REST endpoints for session management, geocoding, chat, preferences, and rankings
- Integrated Gemini 2.5 and Google Maps APIs to rank top **3** restaurants from up to **50** candidates per search
- Containerized Streamlit frontend, FastAPI backend, and Redis cache with Docker Compose, testing **10** simultaneous lobbies

GitHub

Personalized Medicine AI Pipeline | *PyTorch, Python, SQL*

- Engineered PyTorch pipelines for treatment outcome prediction, optimizing model training on sparse clinical datasets
- Deployed low-latency inference service processing high-dimensional patient records under **50ms** for decision support
- Implemented Python imputation workflows for missing clinical data, preserving **99%** integrity for downstream modeling